

Performance Evaluation in Multi-Tenant In-memory Data Grids

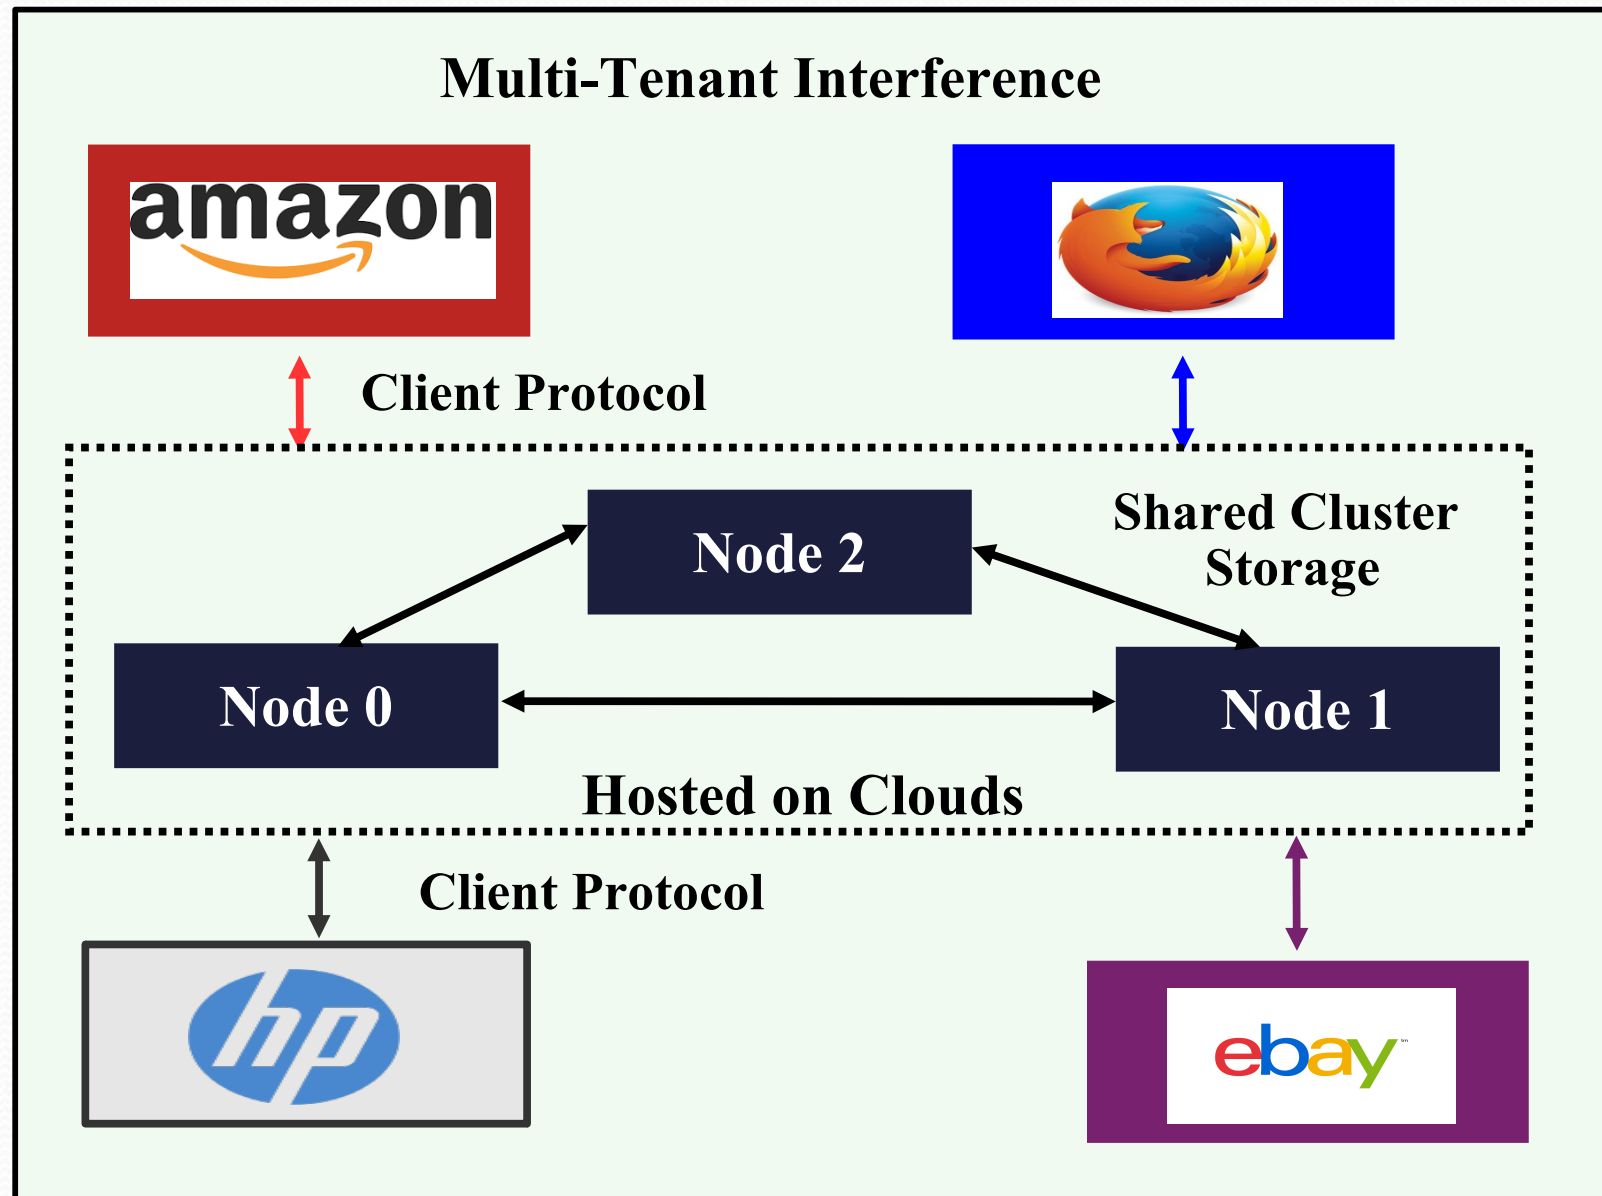
By:
Anwasha Das

North Carolina State University

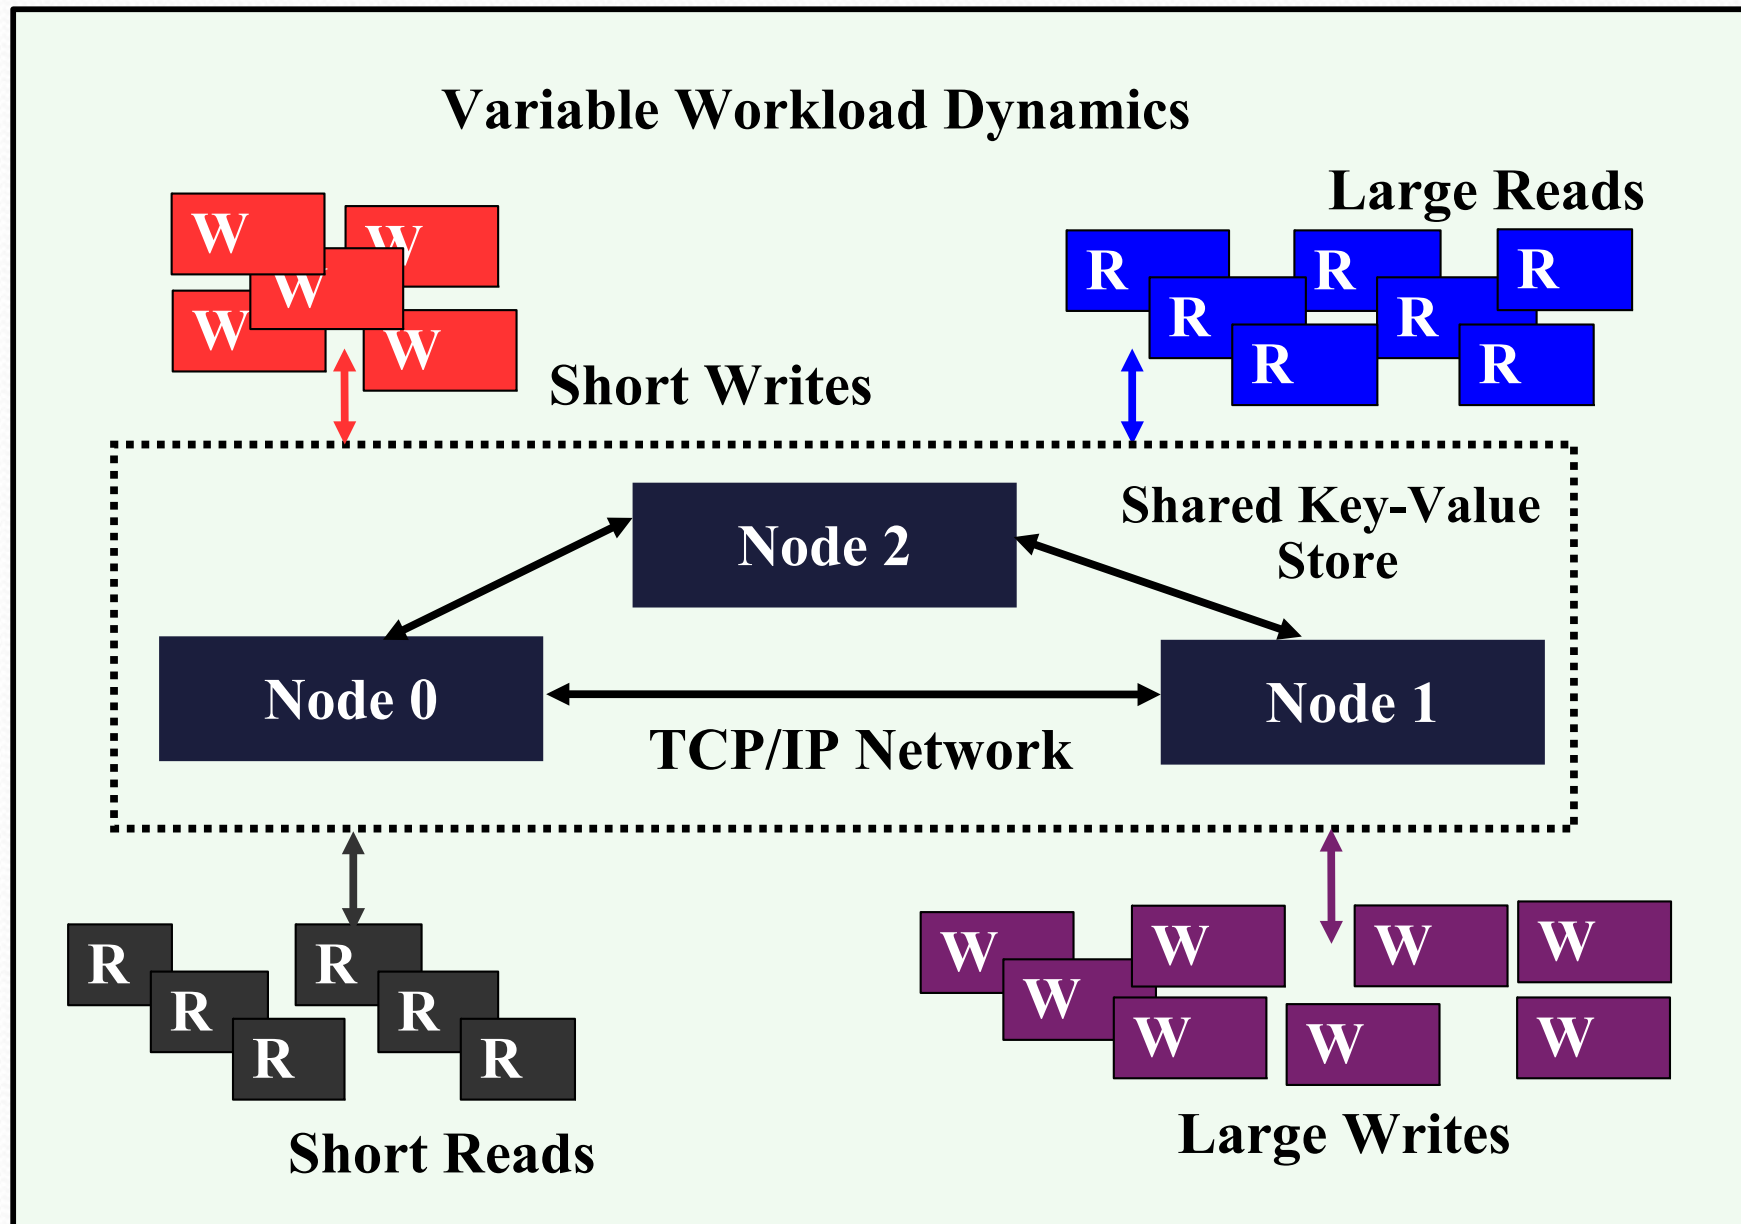
Outline

- Problem
- Research Goals
- Solution Approach
- Experimental Evaluation
- Related Work
- Conclusion/Future Work


Motivation



Motivation




Problem



- ✓ *Decrease in throughput with increase in number of clients*
- ✓ *Presence of contention leading to performance degradation*

Problem



Problems ? Resource Contention, Performance Degradation

Aim – Understand the source of contention, Find out performance bottlenecks in Hazelcast, How to alleviate contention ??

Key-Value Store Users

Key-Value Stores	Open Source
------------------	-------------

<i>BigTable</i>	<i>No</i>
-----------------	-----------

<i>Pnut</i>	<i>No</i>
-------------	-----------

<i>DynamoDB</i>	<i>No</i>
-----------------	-----------

<i>MongoDB</i>	<i>Yes</i>
----------------	------------

<i>Voldemort</i>	<i>Yes</i>
------------------	------------

<i>HBase</i>	<i>Yes</i>
--------------	------------

<i>HvnerTable</i>	<i>Yes</i>
-------------------	------------

<i>ZBase</i>	<i>Yes</i>
--------------	------------

<i>Cassandra</i>	<i>Yes</i>
------------------	------------

<i>Memcached</i>	<i>Yes</i>
------------------	------------

<i>Redis</i>	<i>Yes</i>
--------------	------------

<i>Hazelcast</i>	<i>Yes</i>
------------------	------------



Wide Commercial and Academic Usage

Outline

- Problem
- Research Goals
- Solution Approach
- Experimental Evaluation
- Related Work
- Conclusion/Future Work

Objective

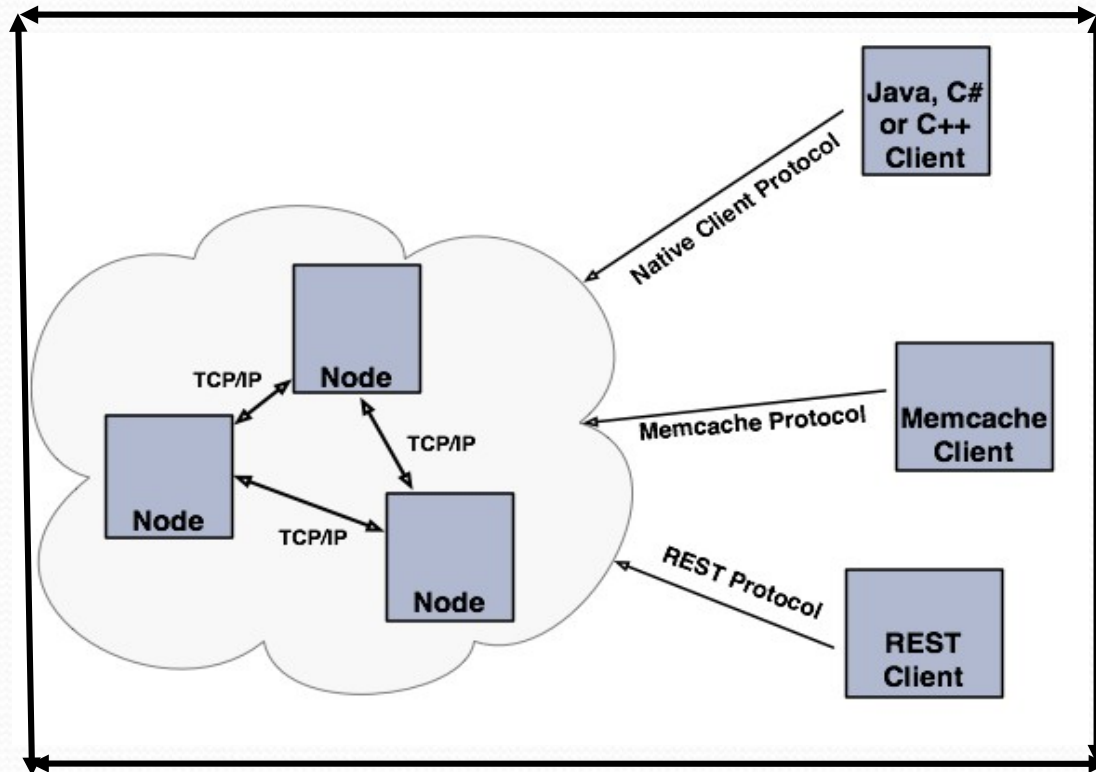
- *Understand Hazelcast in the context of multi-tenancy.*
- *How high is the performance degradation ?*
- *Observe performance characteristics with varying number of clients, workload type, target throughput, thread count.*
- *What causes degradation ? Investigate performance bottlenecks in Hazelcast to eliminate contention.*

Outline

- Problem
- Research Goals
- Solution Approach
- Experimental Evaluation
- Related Work
- Conclusion/Future Work

Hazelcast Architecture


Client to Cluster



Hazelcast

Client – No data
Default – 271 Partitions

Peer to Peer



Node

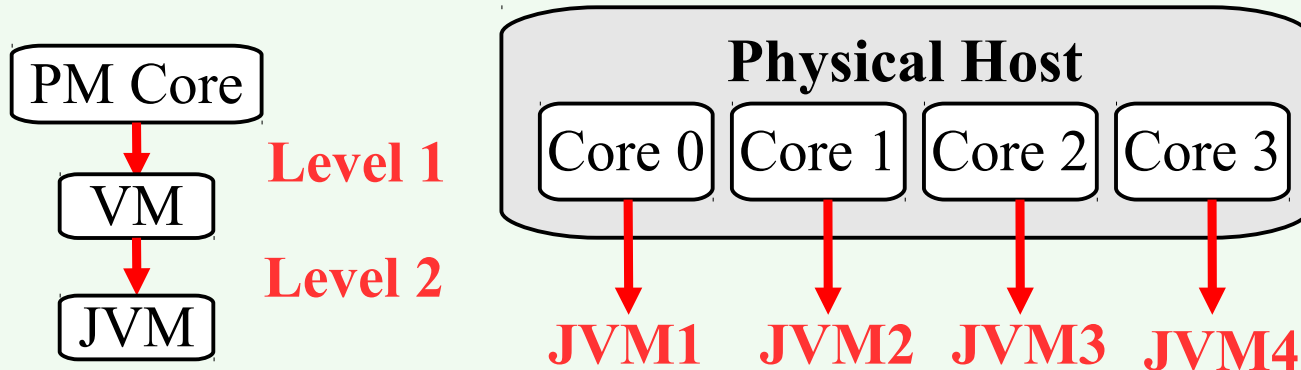
Data

Primary

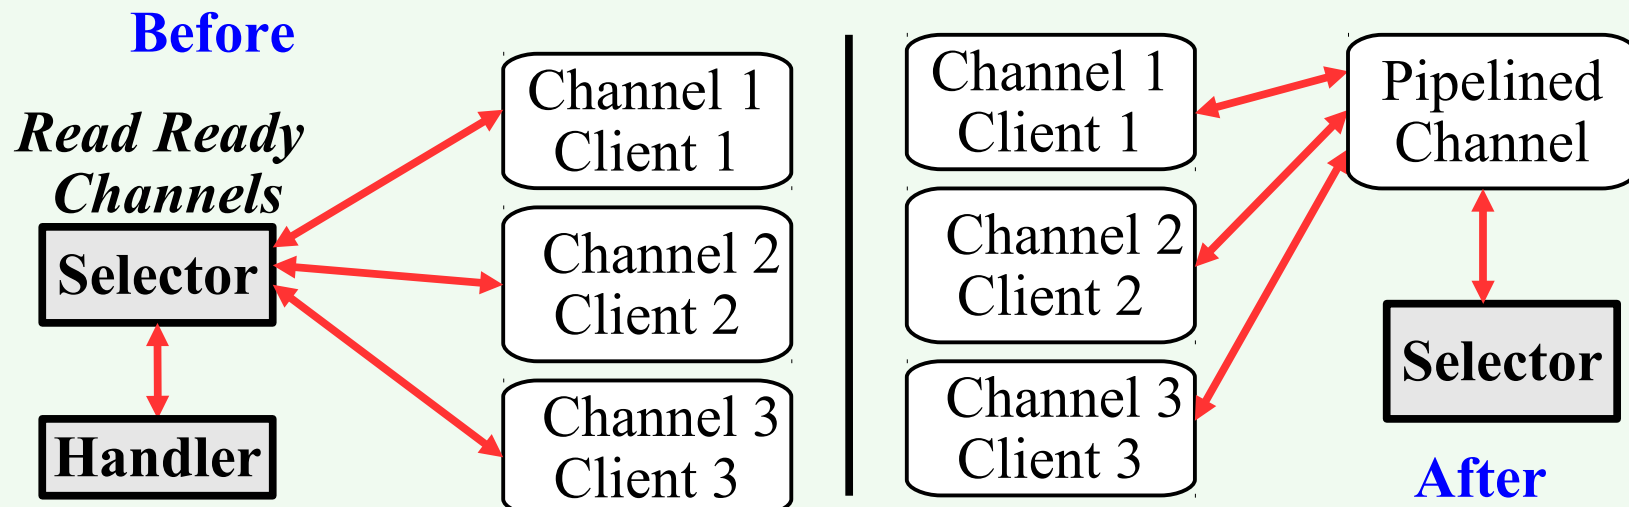
Replica

Solution Approach

Pinning JVMs to CPU Cores



Pipelining of Client Socket Channels



Outline

- Problem
- Research Goals
- Solution Approach
- Experimental Evaluation
- Related Work
- Conclusion/Future Work

Implementation

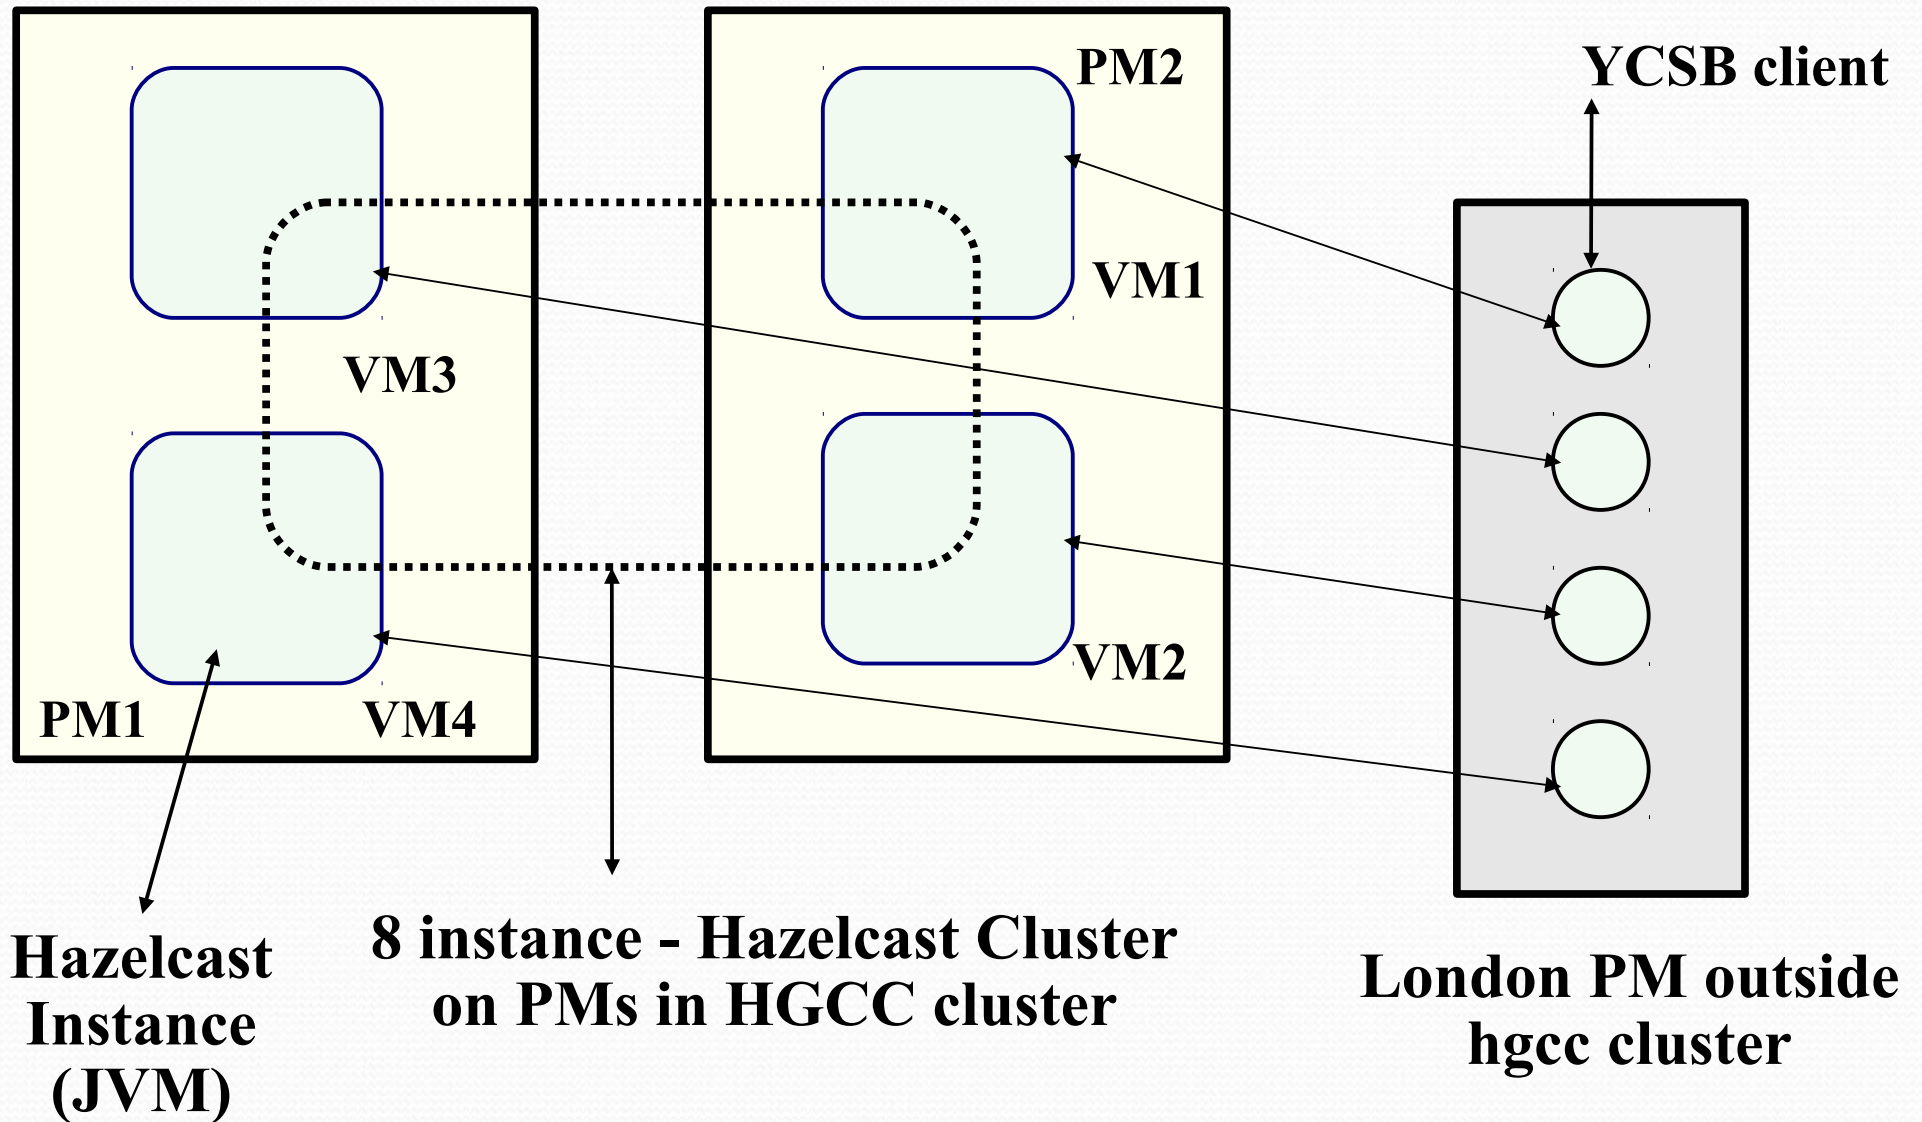
Software

- Java \longleftrightarrow *Hazelcast Code*
- Netty Library Used
- YCSB – Yahoo Cloud Serving Benchmark \longleftrightarrow *Evaluation*

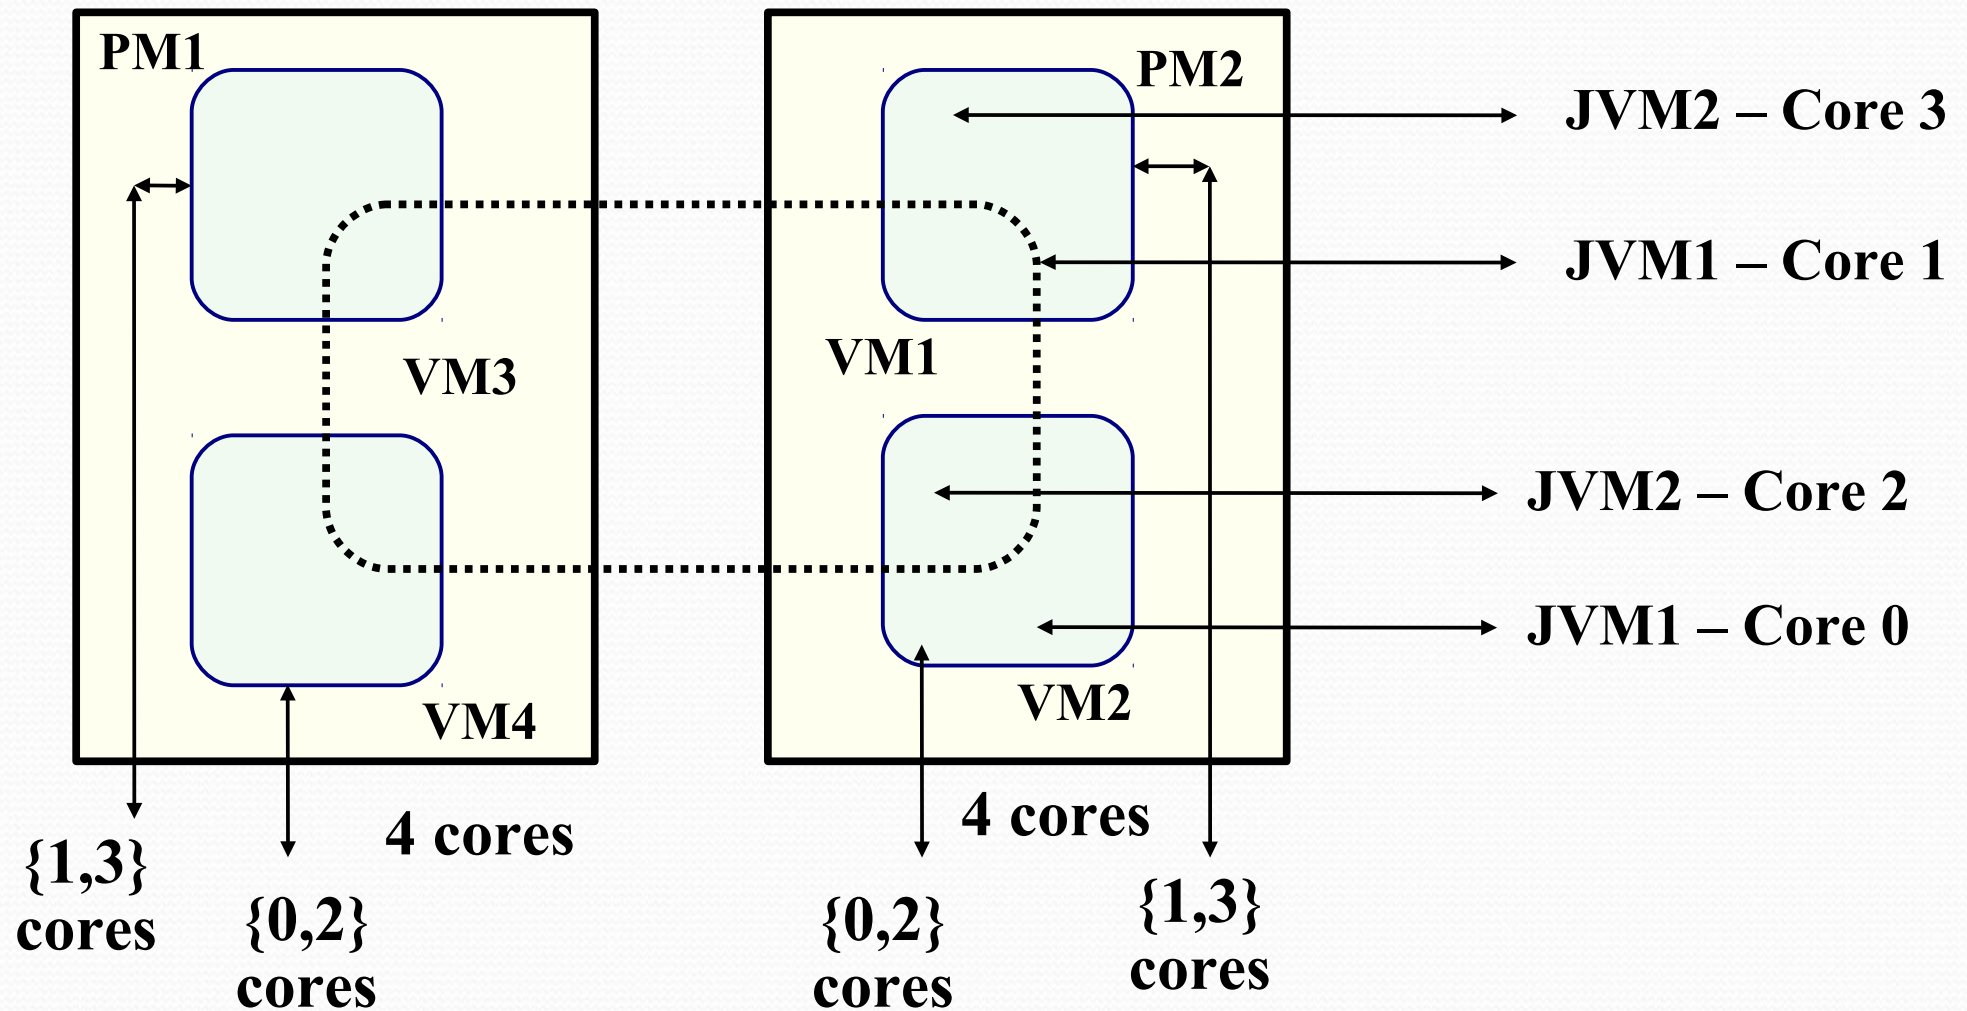
Hardware/OS Platform

- Nodes on HGCC Cluster
 - Quad-core Xeon 2.53GHz CPU, 8GB memory
- VMs running with Ubuntu 12.04 32bit with 4 GB memory and 2 vpcus on 2 HGCC nodes


Experimental Set-Up



Set-Up with Pinning




Client Count Variation




- ✓ *Total number of client side threads used to generate workload fixed to 200, threads per client varies accordingly*
- ✓ *No difference in performance between with and without pinning*

Client Count Variation



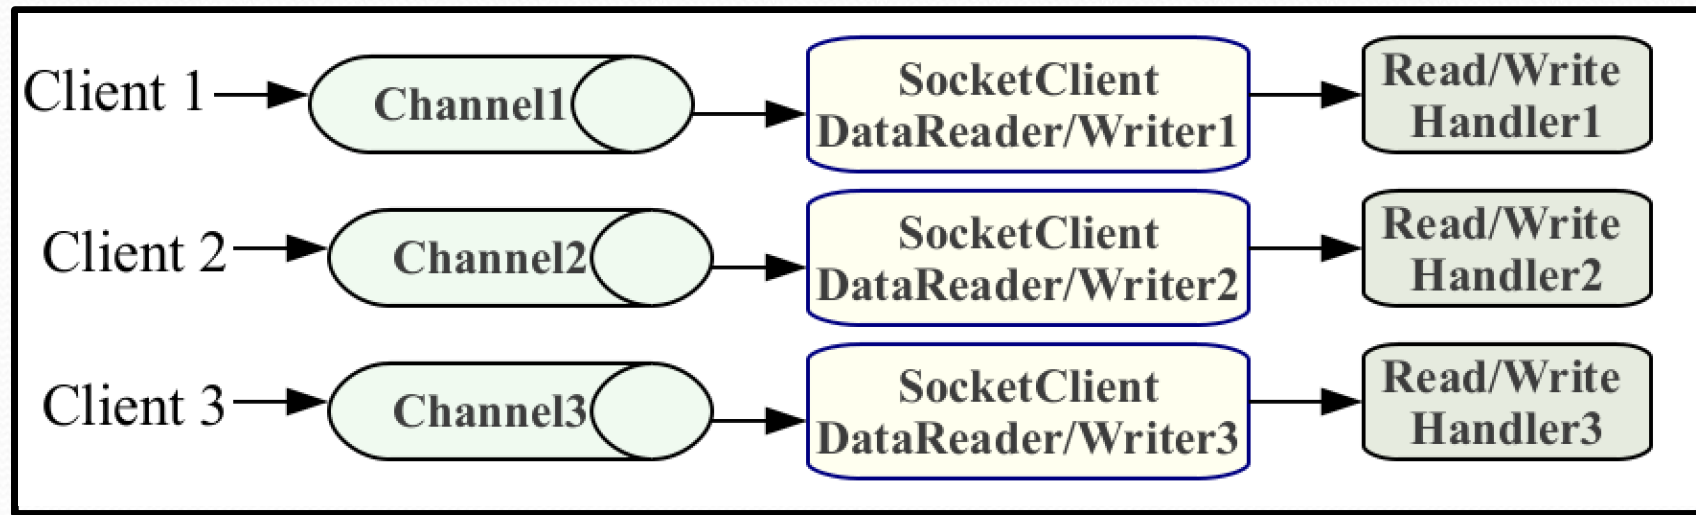
- ✓ *Total number of client side threads used to generate workload increases with every client*
- ✓ *No difference in performance between with and without pinning*

Client Count Variation




- ✓ *Pinning does not help in performance improvement*
- ✓ *Thread migration/Context Switches across cores is not high enough to affect performance*

Multiplexing Socket Channels

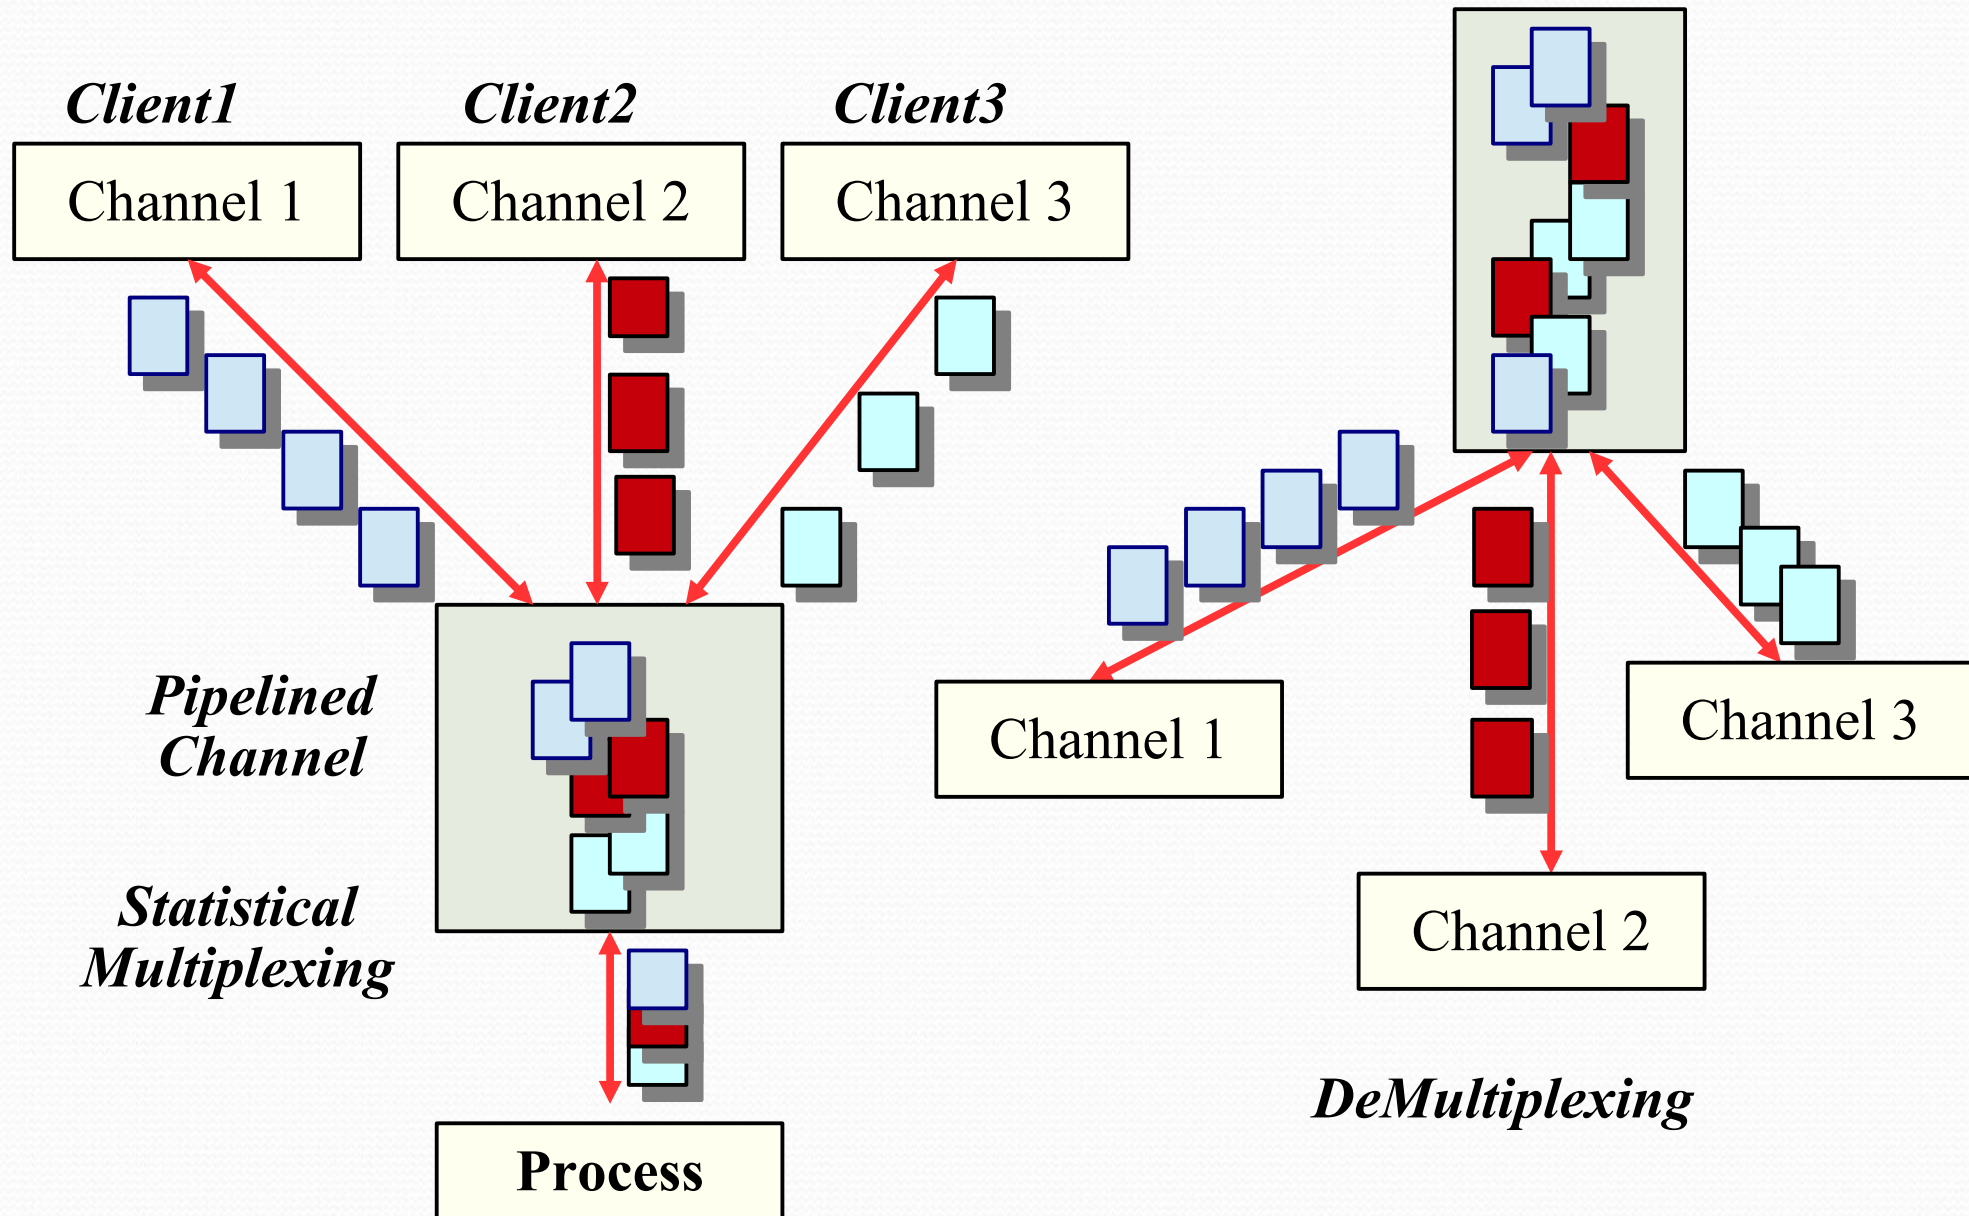


Before




After

Multiplexing Socket Channels




Multiplexing Socket Channels




- ✓ *Performance Improvement*
- ✓ *% increase in throughput as high as 20% for 5 clients*

Multiplexing Socket Channels

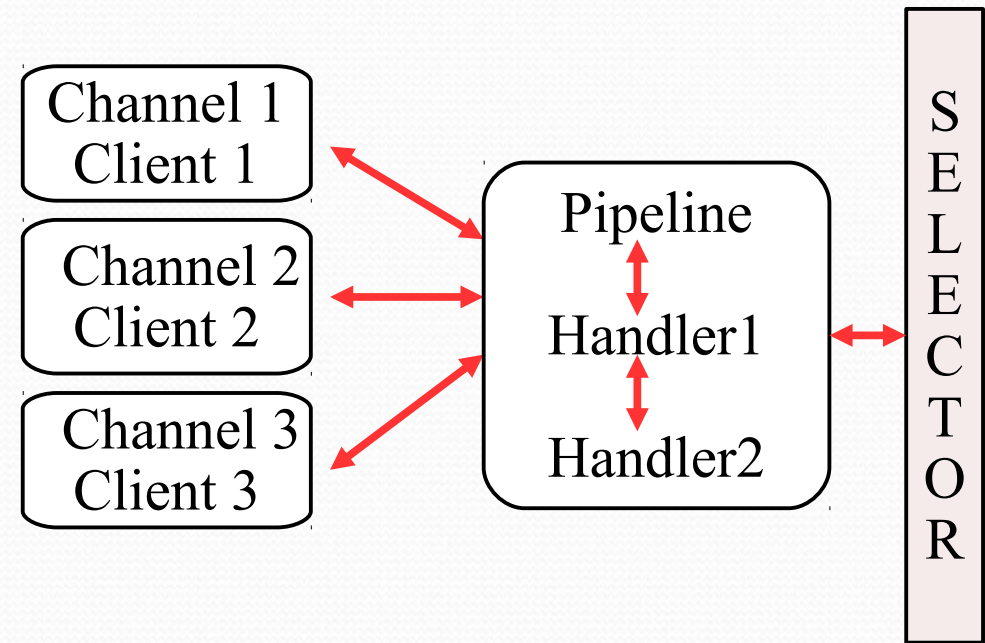


- ✓ *Standard Deviation did not exceed 50*
- ✓ *Less endpoint management, better performance*

Multiplexing Socket Channels

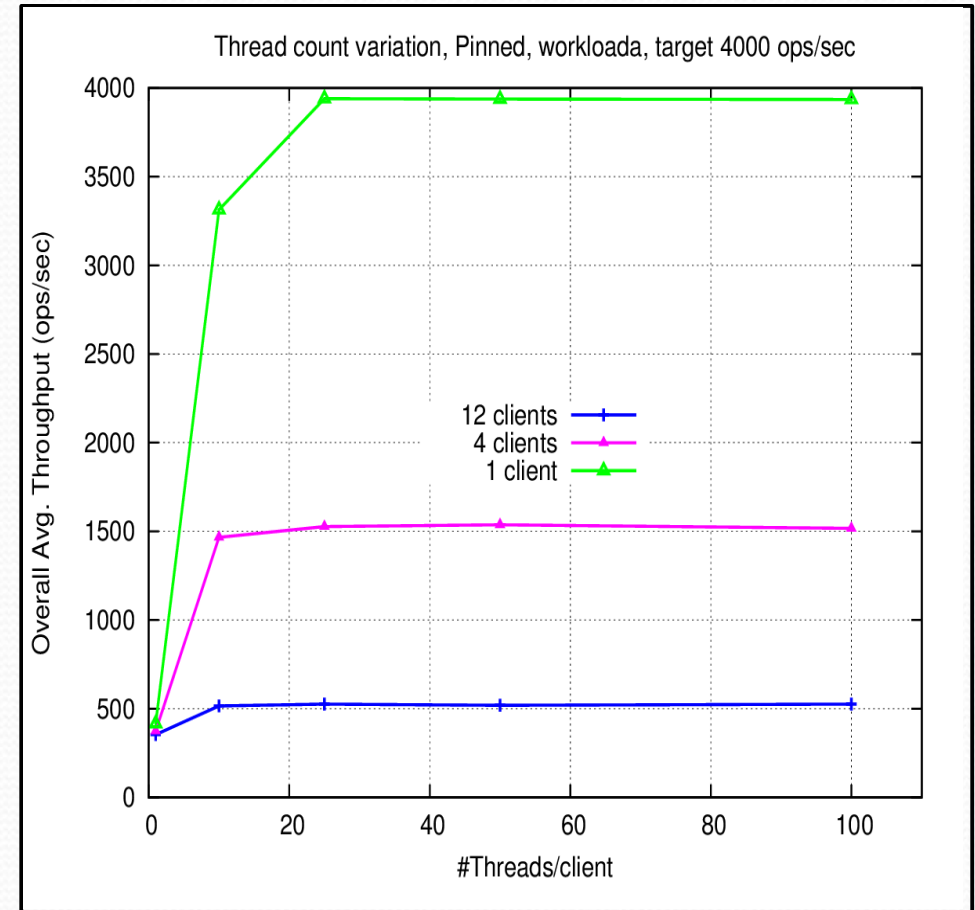
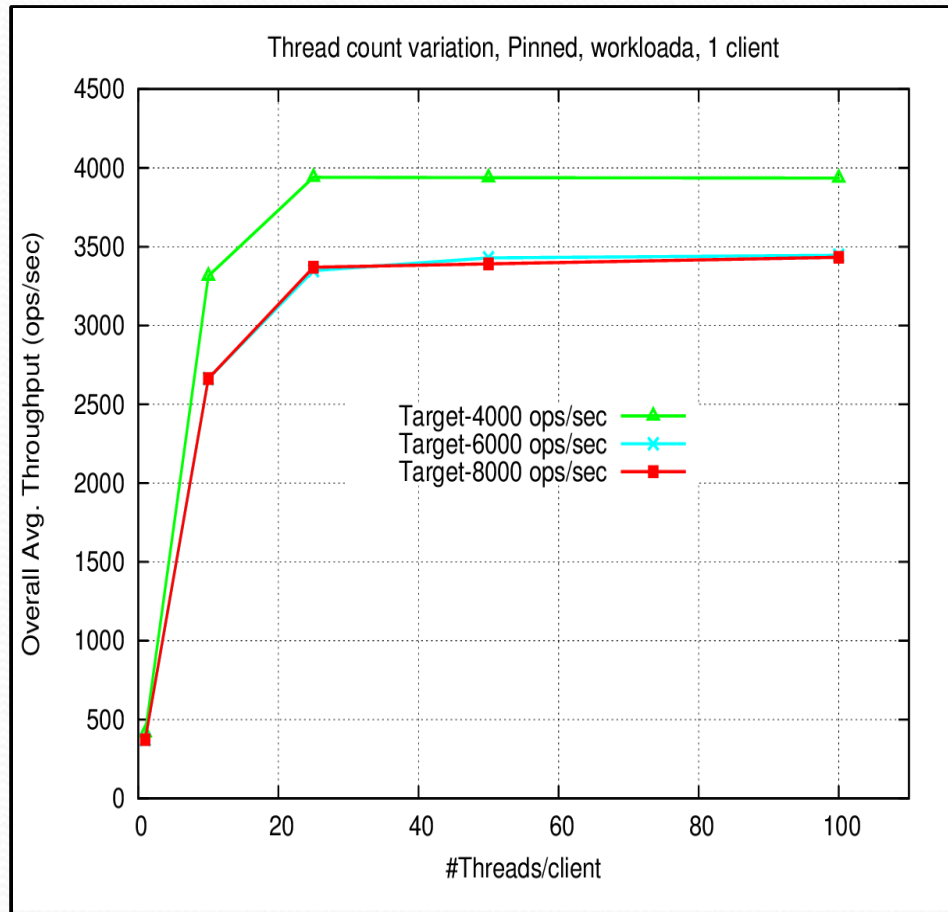


- ✓ *Select/Poll*
- ✓ *More Context Switches*
- ✓ *More Runnable Instances - higher threads per request*



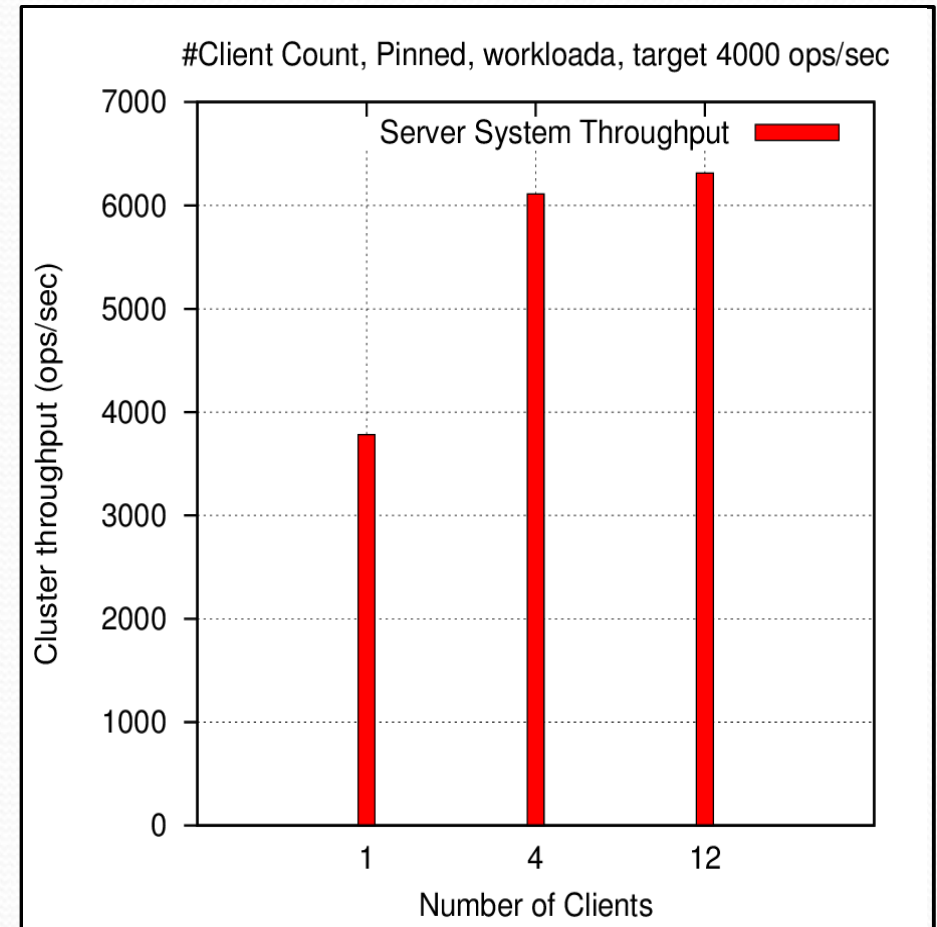

- ✓ *Less Context Switches*
- ✓ *Less threads per request*
- ✓ *Better performance with concurrency*

Client Thread Count Variation



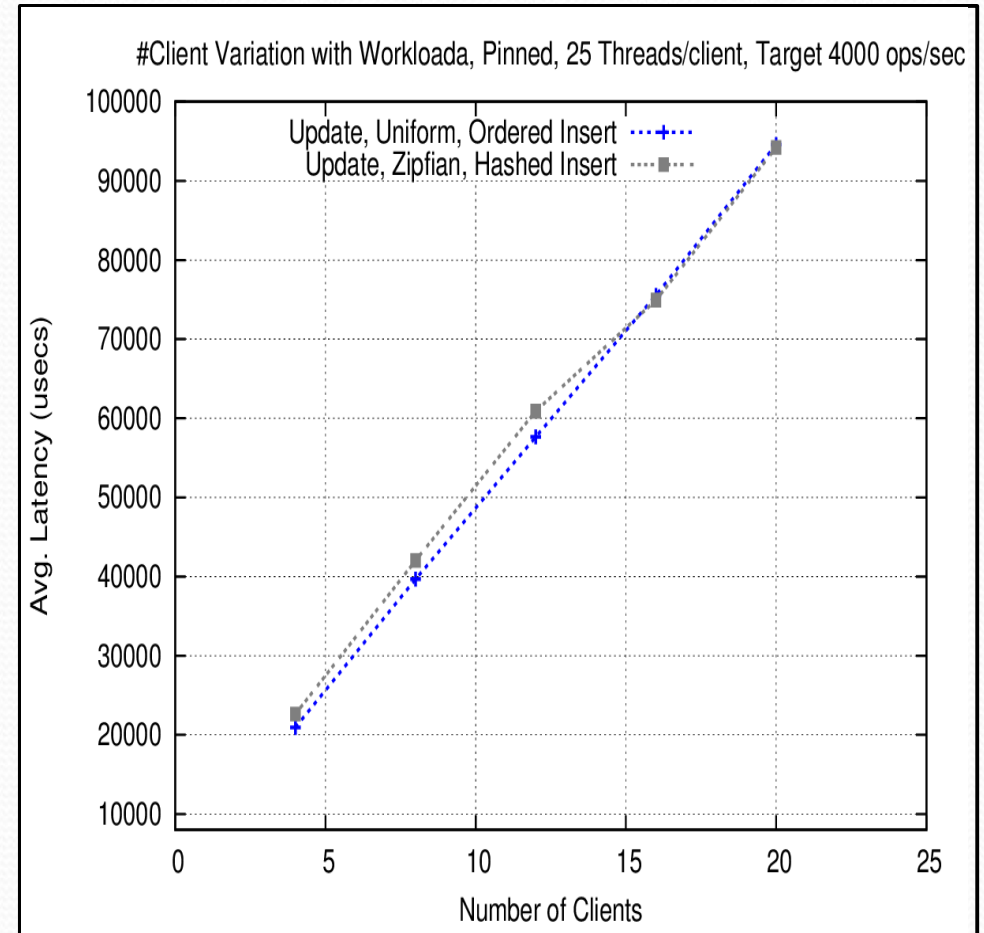
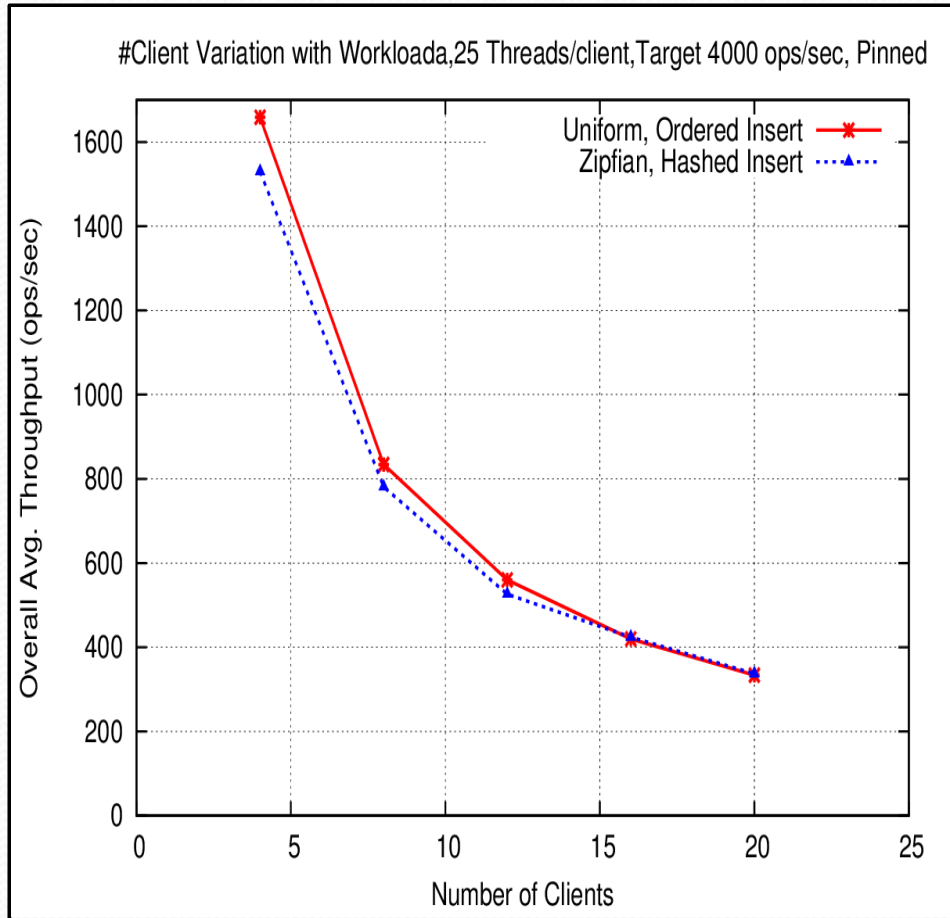
- ✓ *Single client achieves target, beyond 4000 ops/sec performance drops.*
- ✓ *With both 4 & 12 clients, system reaches maximum performance level at 25 threads/client beyond which there is no fluctuation.*

Target Throughput



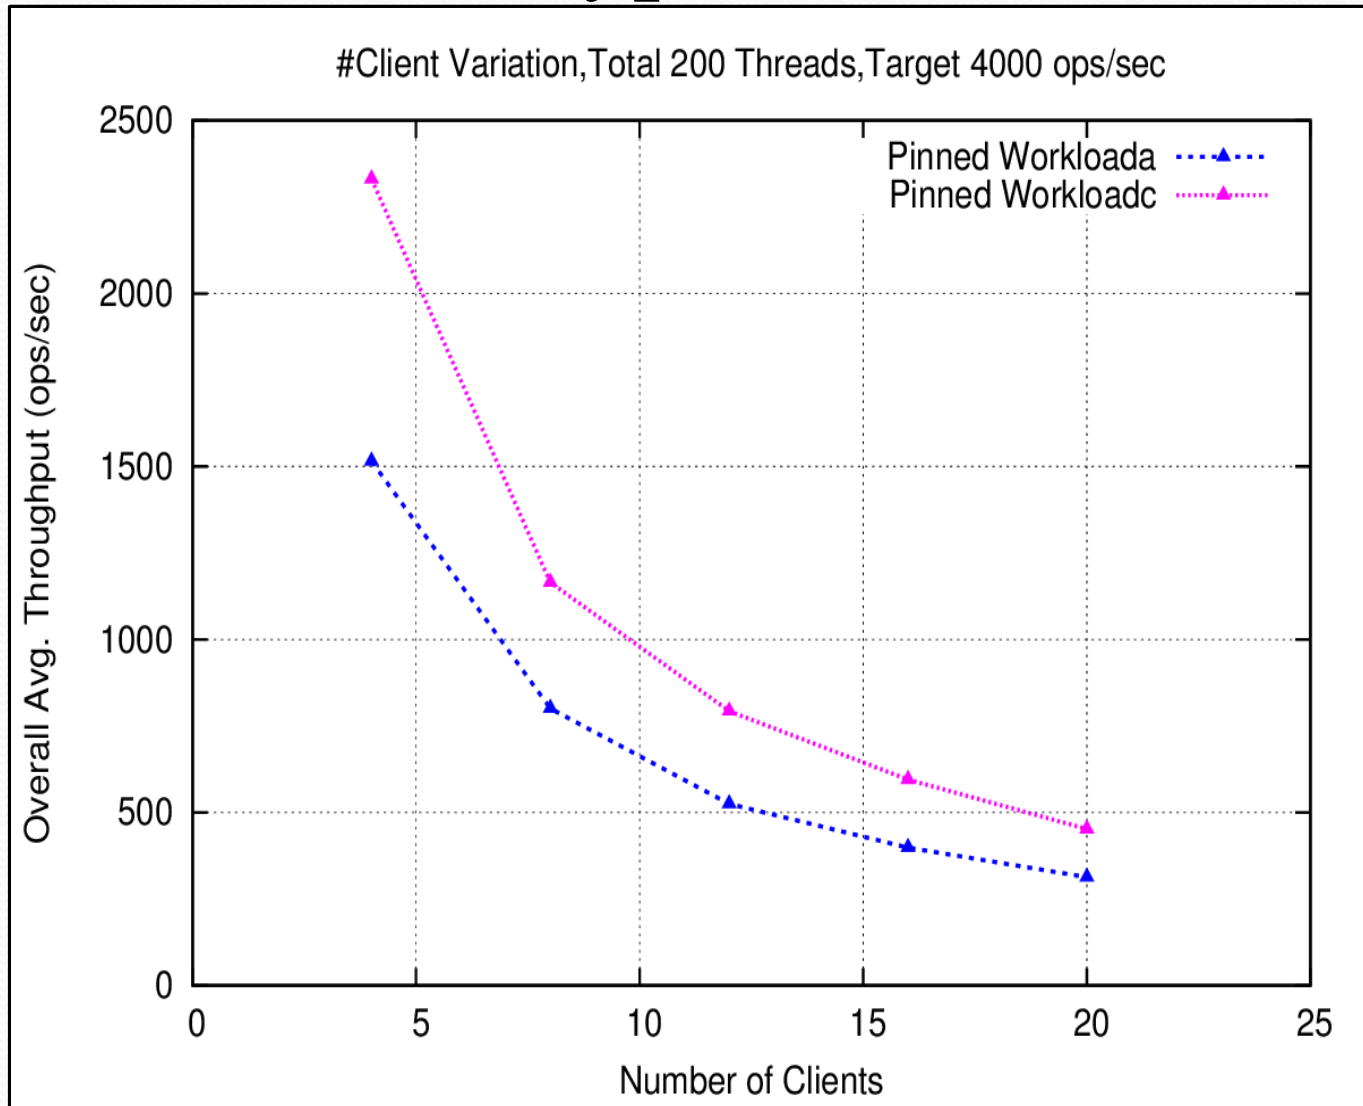
- ✓ *Target doubled, no difference in throughput with the same #clients.*
- ✓ *Cluster throughput from the server's perspective does not exceed 6300 ops/sec.*

Insert Order & Distribution Type



Whether ordered or hashed inserts, irrespective of uniform or zipfian distribution, the performance degradation is similar.

Workload Type Variation




- ✓ **Workloada –**
 - 50% read,**
 - 50% update**
- ✓ **Workloadb –**
 - ✓ **95% read**
 - ✓ **5% update**
- ✓ **Workloadc –**
 - ✓ **100% read**
- ✓ **Workloadd –**
 - ✓ **95% read**
 - ✓ **5% insert**
- ✓ **Workloaddf –**
 - ✓ **50% read**
 - ✓ **50% read-modify-write**

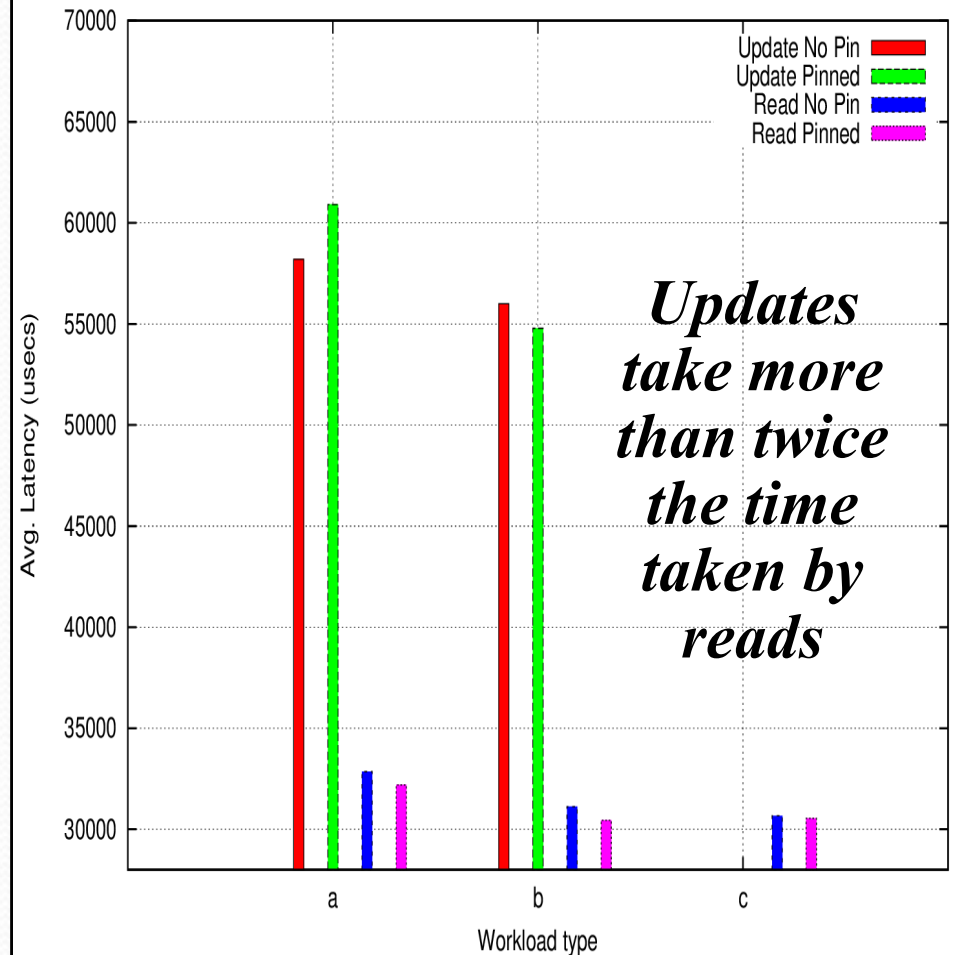
Even read-only workload suffers throttled throughput.

Workload Type Variation

Workload variation with 12 clients, 25 threads/client, target-4000 ops/sec



Workload variation with 12 clients, 25 threads/client, target-4000 ops/sec



More updates, lower the throughput, read-modify-write costlier than update.

Outline

- Problem
- Research Goals
- Solution Approach
- Experimental Evaluation
- Related Work
- Conclusion/Future Work

Related Work

- Centralized Solutions
 - Pisces, Borg, MROrchestrator, HDFS based HFS,
 - Changed scheduling logic
- Decentralized Solutions
 - Apollo, Sparrow, Omega, Mercury
 - Comprehensive scheduler considering diverse heterogenous workloads
- Two-level schedulers
 - Cake, Mesos, Yarn
 - Focus on Hadoop Style Applications
- Cache Replacement Schemes, Cluster Load balancing
 - Gdwheel, Camp, MBal, Scads
- How to process data through nultiple entry points ?
 - Contention elimination in Multi-tenant data grids
 - Another perspective of solving the problem

Outline

- Problem
- Research Goals
- Solution Approach
- Experimental Evaluation
- Related Work
- Conclusion/Future Work

Concluding Remarks



- Performance degradation is clear with increasing clients.
- JVM-CPU Pinning does not help in improving performance.
- Beyond a threshold and a specific target, increasing the number of client threads do not affect performance.
- Insert Order and Distribution Type have no tangible affect at-least if all clients have similar transaction type.
- Statistical multiplexing of client socket channels improve overall throughput.

Future Work

- Investigate efficient ways to process requests through multiple entry points.
- Considering the scale of cluster, can performance be improved through multiplexing, avoiding excessive parallelization ? What are the major limitations ?
- Diagnose *failures* in such systems which affect performance.





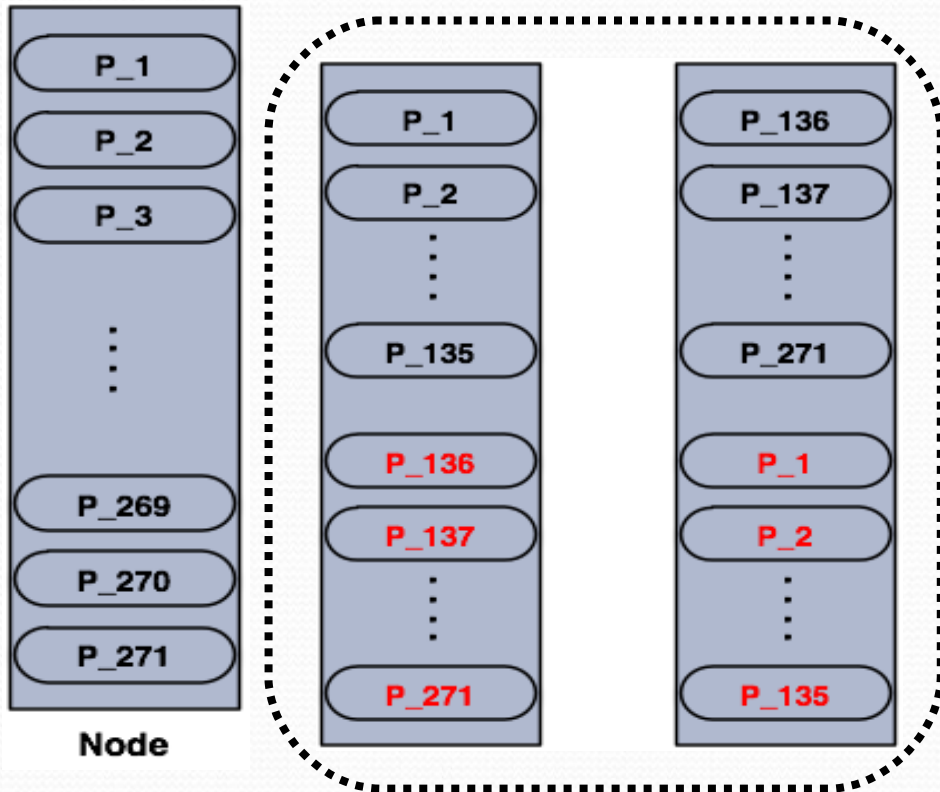
Thank You



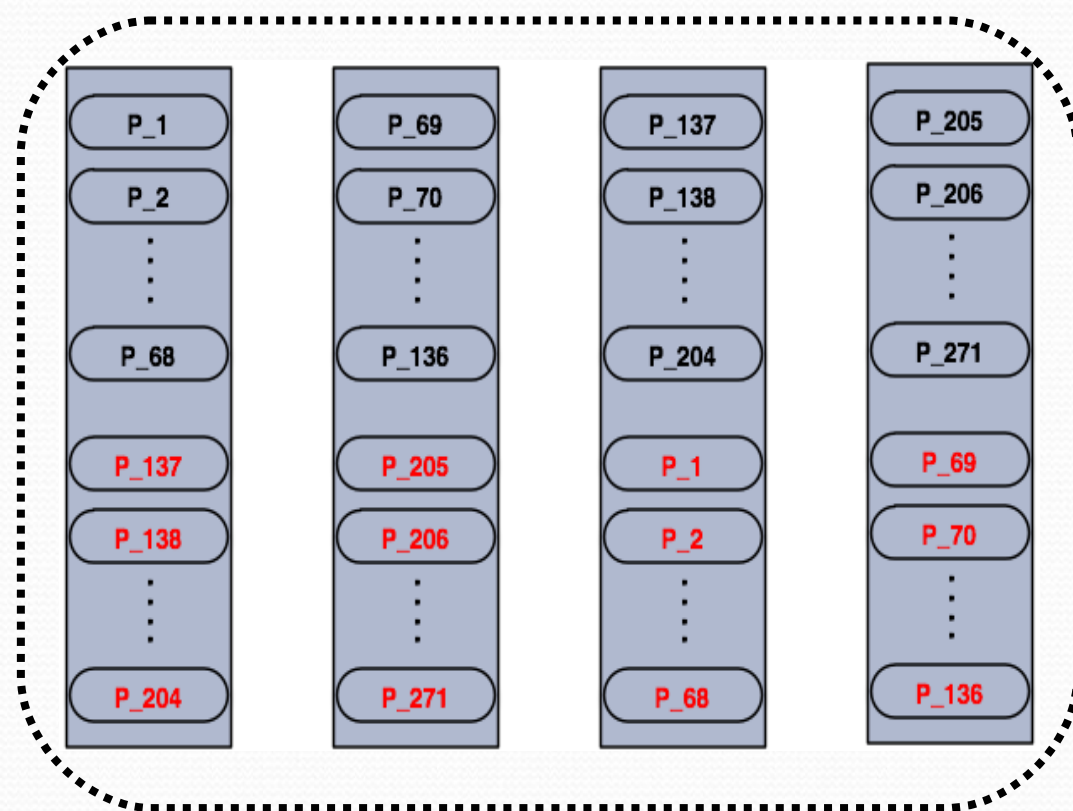
Backup Slides

Architecture

Two Nodes

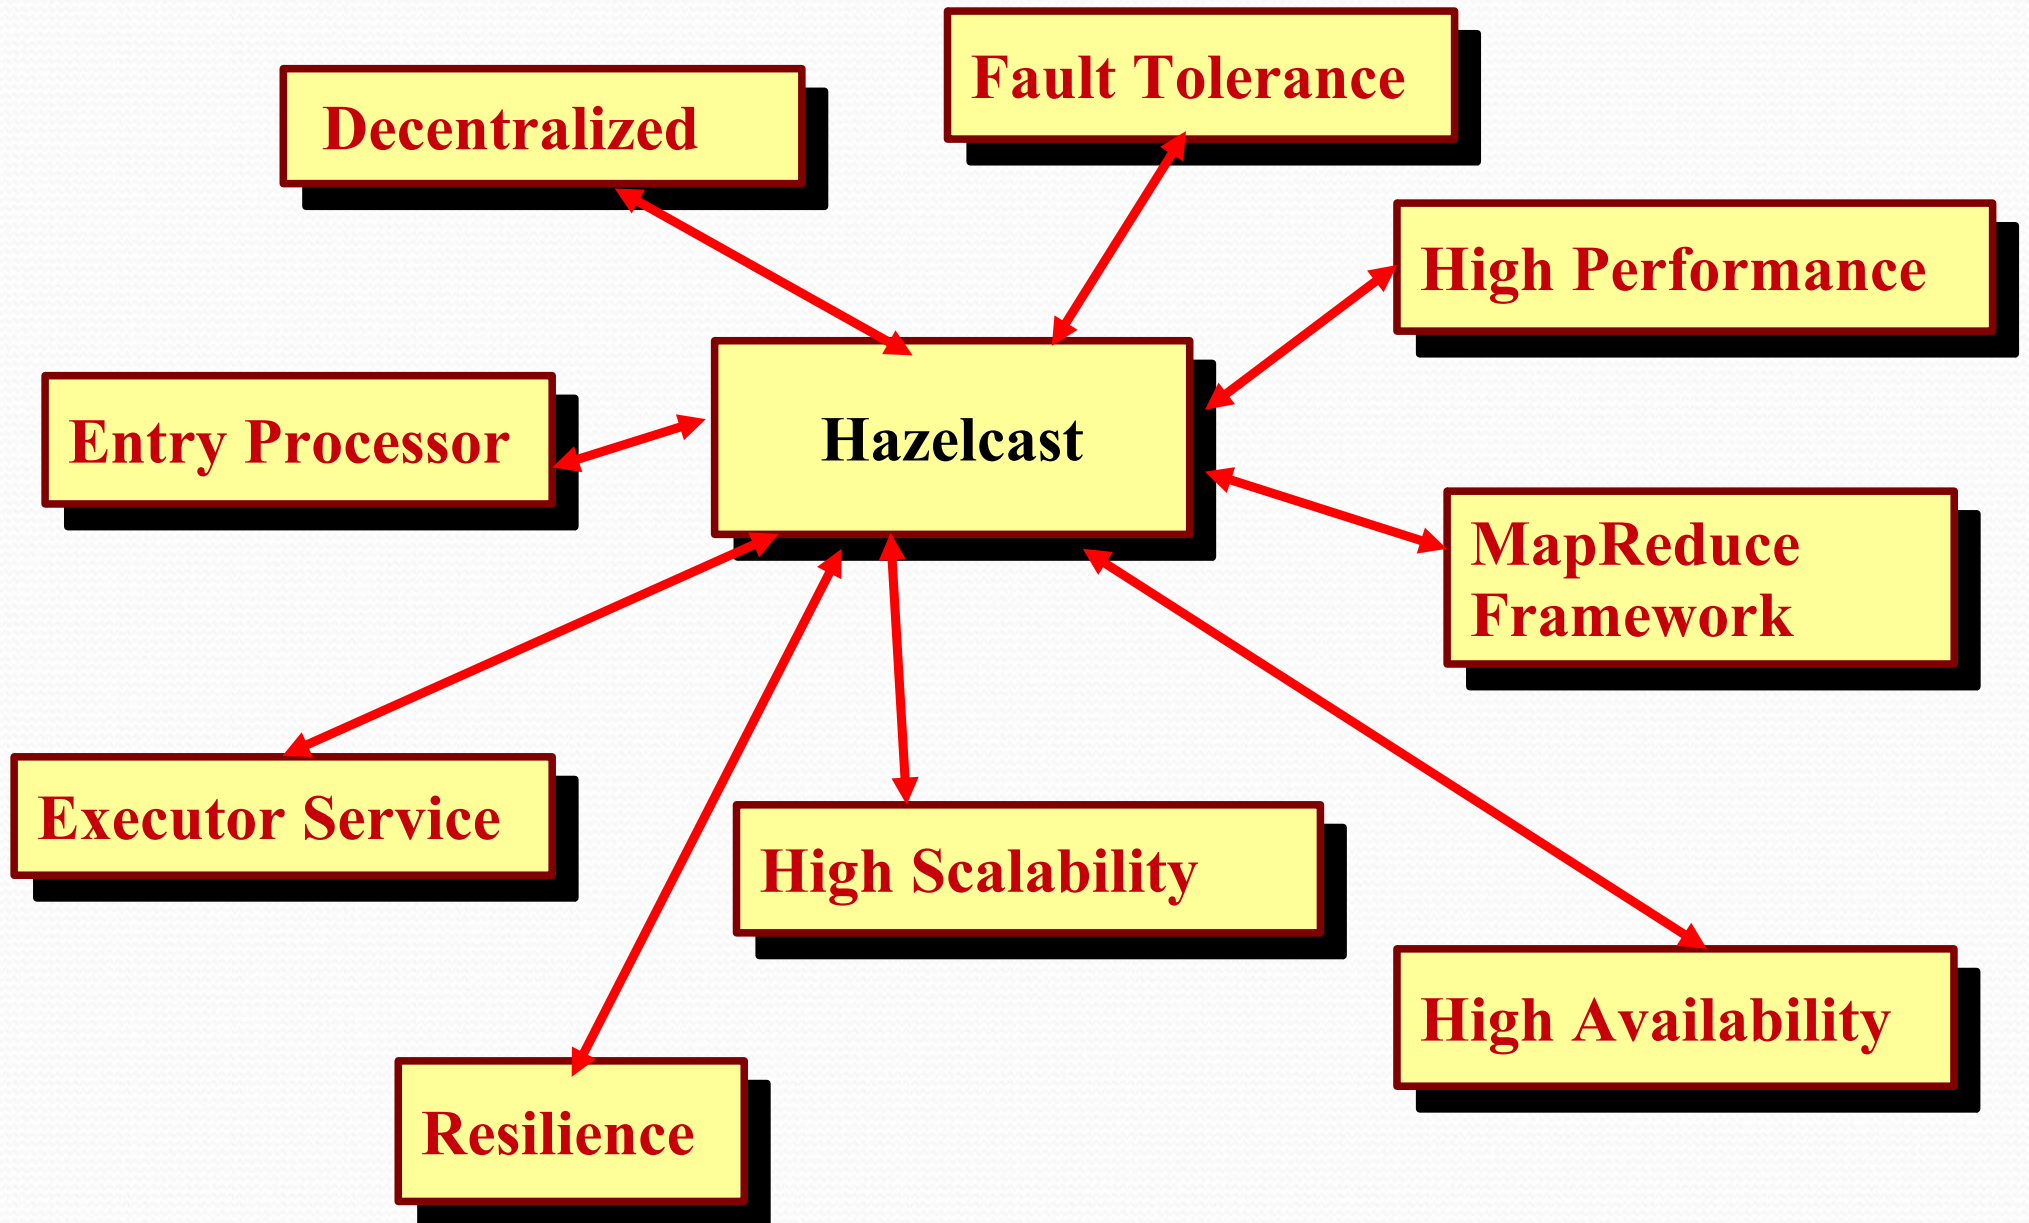


Four Nodes

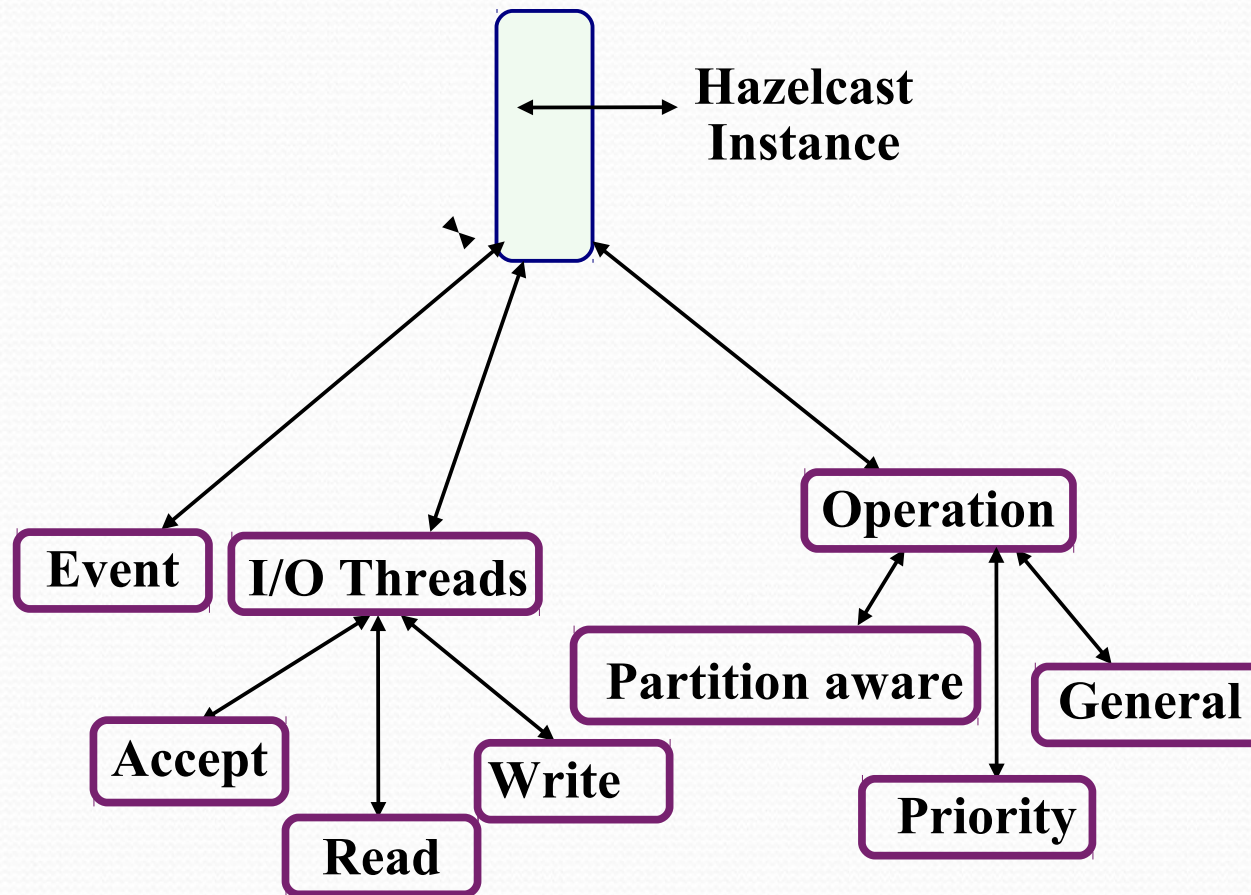


Multiple Instances of Hazelcast cluster in 1 physical Node or across multiple physical nodes possible.

Hazelcast Features



Internal Threading Model

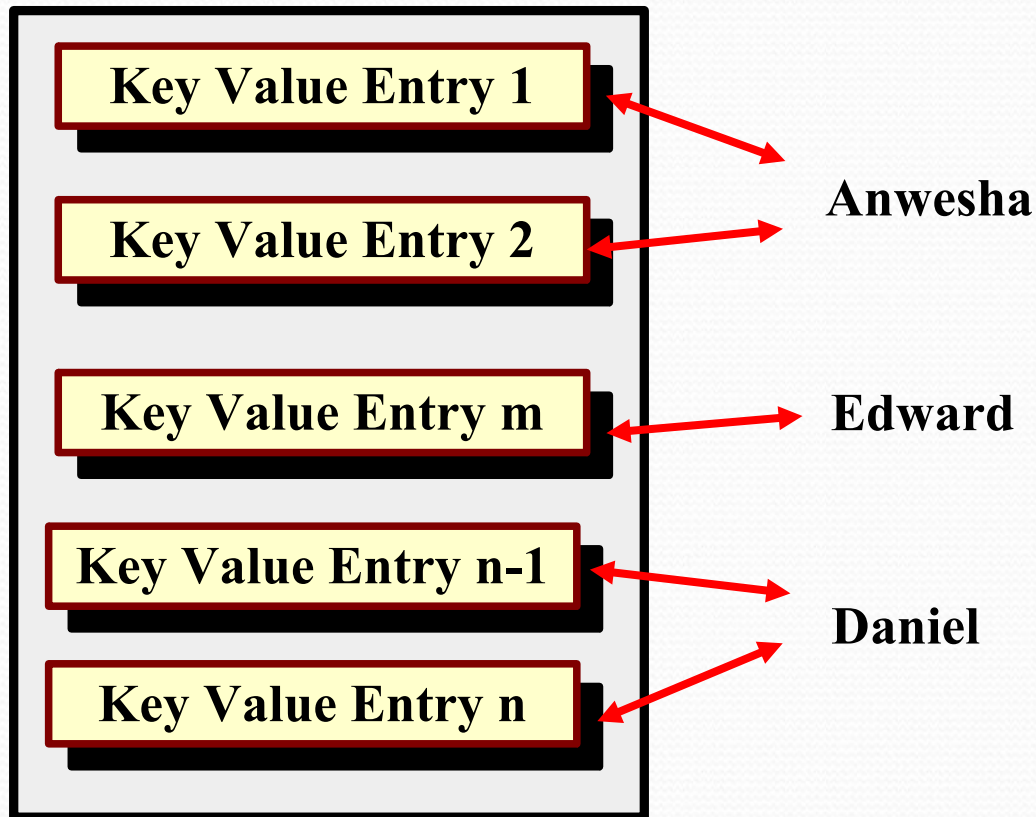


Each thread has its own queue.

Default counts:

- ✓ **Event Threads – 5**
- ✓ **Work Queue capacity - 1000000**
- ✓ **I/O Threads – 7**
 - ✓ **1 accept**
 - ✓ **3 read**
 - ✓ **3 write**
- ✓ **Operation Threads - CPU core count * 2**
- ✓ **Priority Queue**
- ✓ **Iexecutor Threads -**
- ✓ **Executor Service – 16**
- ✓ **Query operator -**
- ✓ **ExecutorPoolSize – 5 * hzClient CPU core count**

Java Client Protocol




**Shared Distributed Map A
across multiple clients**

- ✓ Open Connection
- ✓ Tenant Authentication
- ✓ Retrieve Partition List
- ✓ Send Operation Messages
- ✓ Receive Responses
- ✓ Cluster Member Updates
- ✓ Close Connection


With Single Client (No Contention)

#Client Variation with Workload, 25 Threads/client, No Pinning




With even higher targets, multi-tenancy will give no better performance

Target Perf. Variation, Pinned, 1 client, 25 threads, Workload





Threshold – 4000 ops/sec

With Single Client (No Contention)



High Drop from 1 client to 4 client, Read-only better throughput !!

Client Count Variation



- ✓ *Well distributed workload, better overall response time*
- ✓ *Excessive parallelization w.r.t target increases latency*